


Concurrency and parallelism in Rust



Why concurrency?


- Sometimes things need to happen at the same time
 - Computer inputs and outputs cannot always wait
 - Sometimes you wait for external resources to finish, e.g. I/O
 -
-

CPUs and concurrency

- CPUs come in many flavours:
 - Single-core CPUs permits a single active thread of execution, i.e. *one thing at a time*

Multi-cores and parallelism

- Having more than one CPU-core allows executing more than one thread of instructions at any time
- *True* parallelism
- Different naming between OS's (naming is hard), but generally called *processes*, which can spawn *threads*, which themselves can be run on different CPU cores, dependant on OS implementation



Parallelism outside the CPU

- Computers rarely work alone
- Multi-CPU setups
- E.g. two CPUs with single-cores each, or many CPUs with many cores each
- Supercomputers, networks, distributed systems, clusters, multiple servers, etc.

Dangers of concurrency

Data hazards

- Non-*atomic* actions on shared data are unsafe to perform in a concurrent context
- Read-after-write: **A** tries to read from store before **B** finishes writing to it
- Write-after-read: **A** reads from store, update data, **B** writes to store, **A** writes, then data from **B** is overwritten
- Write-after-write: **A** has updated data, writes it to store, but **B** wrote before **A** got the chance

Runtime solutions

- Locking mechanisms: mutexes, semaphores, barriers, condvars, critical sections, yield-locking, spin-locking, etc.
- Pros and cons
 - 👍 Simple, tons of easy-to-access resources
 - 👎 Can lead to deadlocks, livelocks, convoying, resource starvation, indeterminate results, priority inversion, instability, etc.
 - 👎 Runtime overhead, possibly a lot more than a non-negligible amount

Compile-time solutions

- Non-blocking algorithms: Failure or suspension of any thread cannot cause failure or suspension of another thread
 - Usually implemented through atomic CPU instructions such as compare-and-swap and clever algorithm design
 - Can be wait-free, lock-free or obstruction-free, ordered from greatest to weakest guarantees
 - Pros and cons:
 - 👍 No runtime-overhead of lock synchronization
 - 👍 Fewer hazardous pitfalls
 - 👎 Very difficult to implement and may not always be possible

Other solutions

- Formal verification techniques
 - Process calculus, e.g. π -calculus, CSP (communicating sequential processes)
 - TLA+, language for designing, modelling and verifying distributed and concurrent systems
 - Pros and cons:
 - 👍 Formally verifies the absence of bugs such as race conditions
 - 👎 Requires advanced and niche expertise
 - 👎 Often (e.g. TLA+) requires you to model your program, verify it, and keep the model in sync with the implementation

Other compile-time solutions

- Ownership and lifetimes
 - Rust's ownership and lifetime system guarantees safe-to-execute code in concurrent systems
 - Pros and cons:
 - 👍 Successful compilation guarantees data race-freedom
 - 👍 Simpler to program with, alerted of bugs early
 - 👎 Doesn't guarantee race condition freedom (which only formal verification does)

How does Rust do it?

- `T` means you own a variable of type `T`
- `&'a T` means you borrow an immutable reference to `T` with lifetime `'a`
- `&'a mut T` means you borrow an exclusive reference to `T` with lifetime `'a`
- `T: Send` means `T` can be safely sent across thread boundaries. Auto-implemented if available
- `T: Sync` means `T` can be safely synchronized/shared between threads

What are **Send** and **Sync** types?

Examples:

- `bool: Send + Sync`
- `AtomicBool: Send + Sync`
- `&T: !Send + !Sync`
- `Mutex<T>: Send + Sync where T : Send`

Compiler figures it out based on recursively checking contained types for **Send** and **Sync**-bounds or non-bounds.

Threads in Rust

- Spawning a thread:

```
fn spawn<F, T>(f: F) -> JoinHandle<T> where  
    F: FnOnce() -> T + Send + 'static,  
    T: Send + 'static
```

- Can be read as:
 - Closure **f** returns a type **T** and **f** must be sendable across a thread boundary (e.g. cannot reference local unsendable data), must live for **'static** since it's undecidable when thread dies.
 - **T** must be sendable to another thread (since thread returns it)

What about non-static data?

Rust 1.63.0 to the rescue:

```
fn scope<'env, F, T>(f: F) -> T
where
    F: for<'scope> FnOnce(&'scope Scope<'scope, 'env>) -> T,
```

with `Scope`-definition:

```
fn spawn<F, T>(&'scope self, f: F) -> ScopedJoinHandle<'scope, T>
where
    F: FnOnce() -> T + Send + 'scope,
    T: Send + 'scope,
```

Scoped threads usage

```
let mut a = vec![1, 2, 3]; let mut x = 0;
std::thread::scope(|s| {
    s.spawn(|| {
        println!("hello from the first scoped thread");
        // We can borrow `a` here.
        dbg!(&a);
    });
    s.spawn(|| {
        println!("hello from the second scoped thread");
        // We can even mutably borrow `x` here, because no other threads are using it.
        x += a[0] + a[2];
    });
    println!("hello from the main thread");
});
// After the scope, we can modify and access our variables again:
a.push(4);
assert_eq!(x, a.len());
```

Rust doesn't always save you!

You can still mess up, e.g. by deadlocking, etc.

However, Rust markets itself with *fearless concurrency*:

<https://doc.rust-lang.org/book/ch16-00-concurrency.html>